

Kubernetes




Assoc. Prof. Panche Ribarski, PhD | Assoc. Prof. Milos Jovanovik, PhD
Continuous Integration and Delivery 2023

Services

- Pods get updated, destroyed or just plain fail
- Pods are immutable
 - New Pods will be created in their place
 - That means new IP
- Pod DNS is tied to the unique Pod name
 - who would keep a list of running Pods? Not us
- Enter Services
 - stable IP
 - stable DNS
 - stable ports



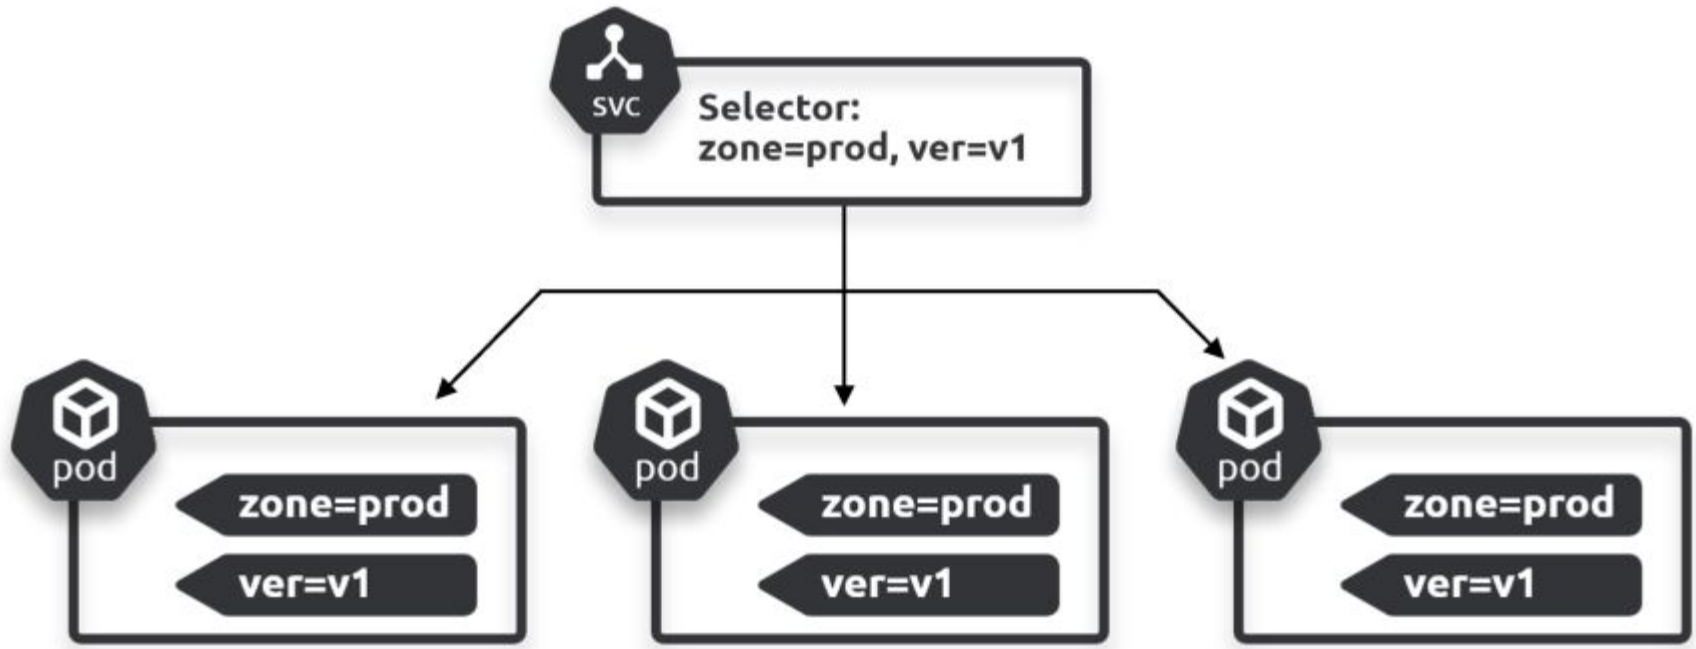
Service object in Kubernetes



Service benefits

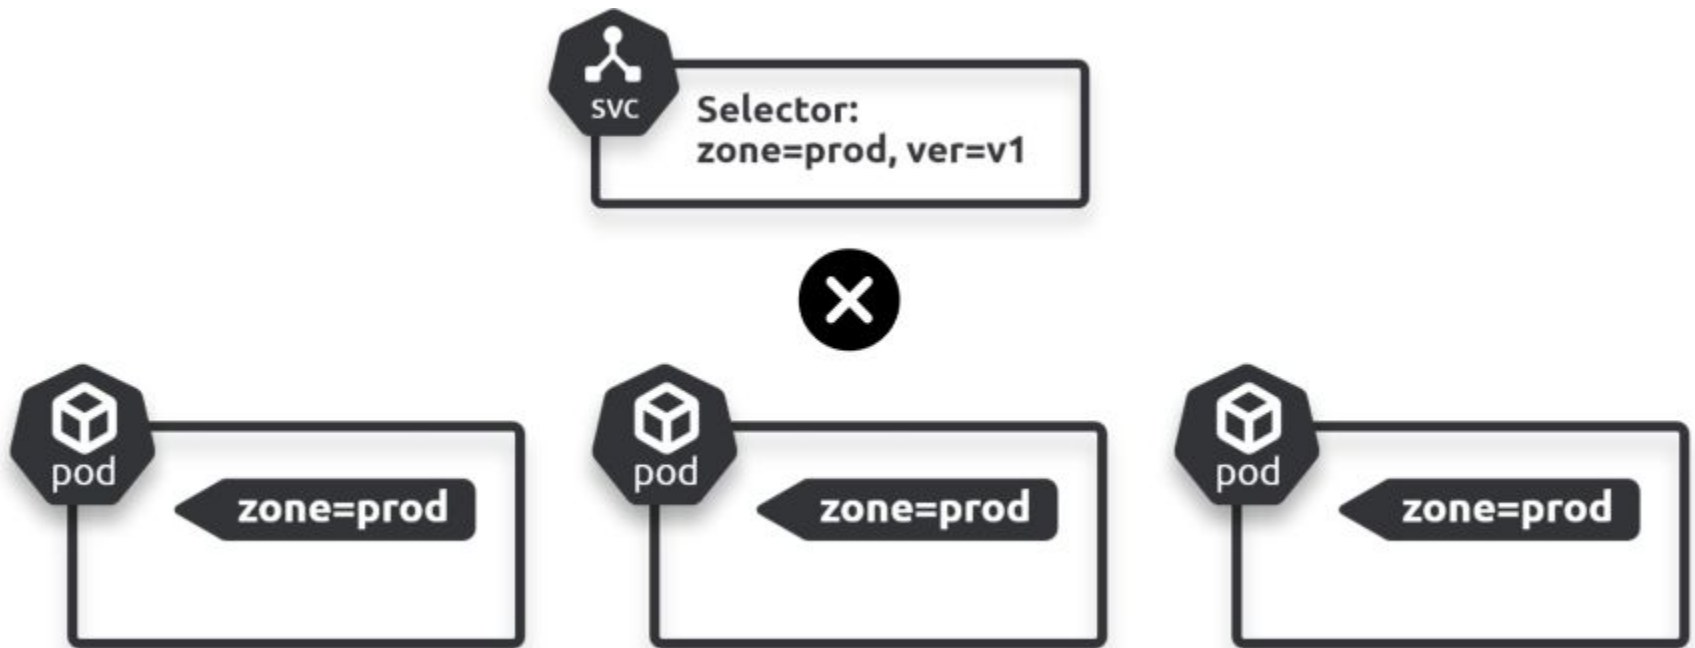
- Deployments can overlook Pods
 - scale
 - update (Rollout)
 - revert (Rollback)
- Everyone else need a stable handle that links to those Pods from the Deployment
- Create a Service object
 - will point to all the Pods (even LB them)
 - keep a list of (healthy) Pods
 - and their IP's
 - index Pods via
 - labels
 - selectors

Service coupling to Pods



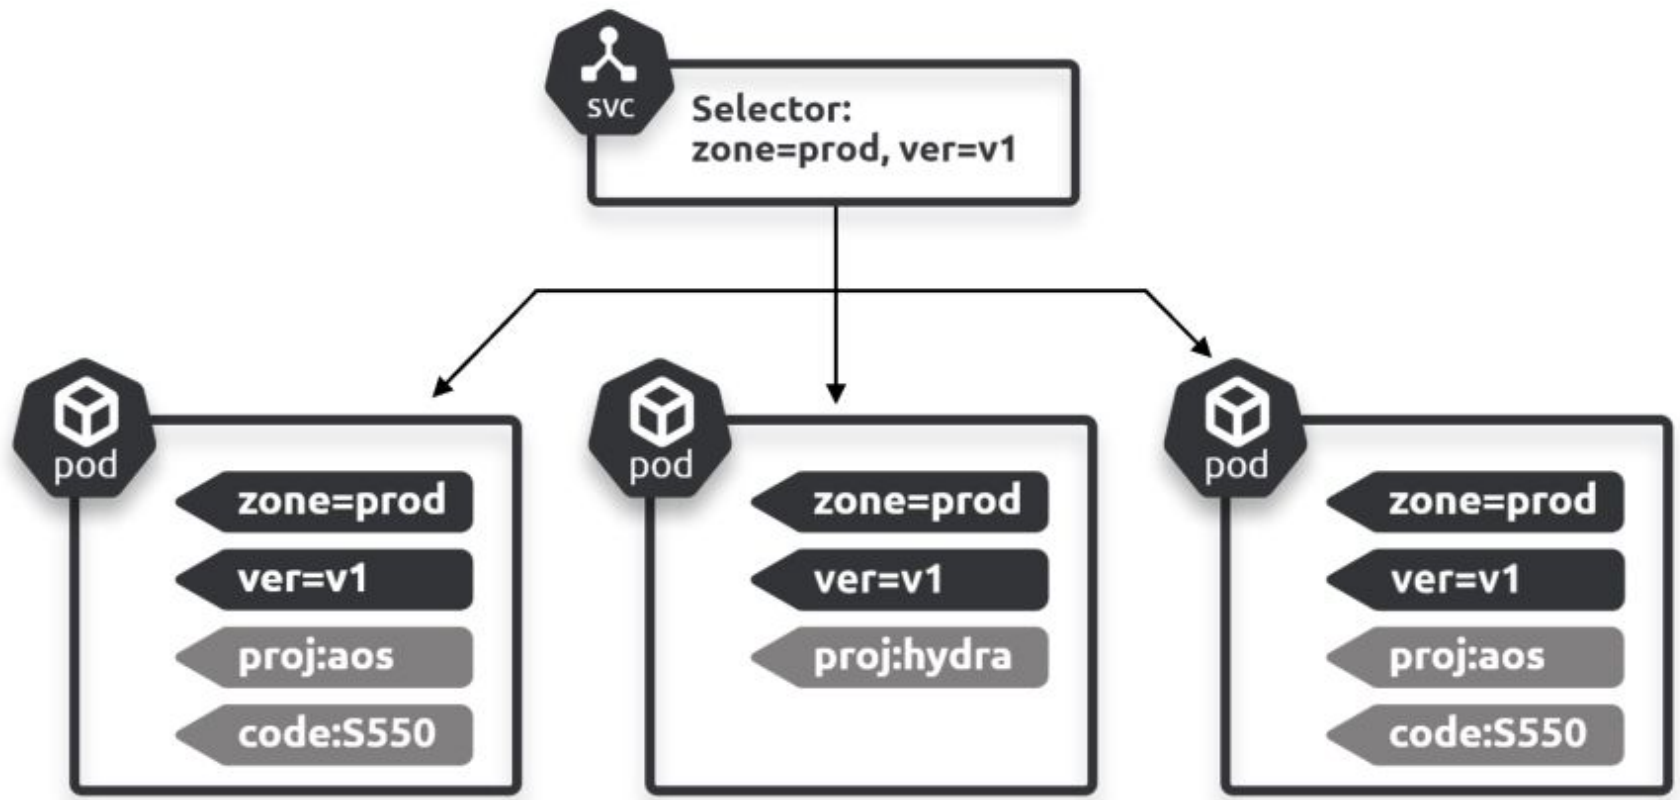
Service selectors

- Service will look for Pods with the listed selectors:
 - all listed tags in selectors must be labels in the Pods



Service selectors

- Your Pods can have extra labels



Our first Service

```
apiVersion: v1
kind: Service
metadata:
  name: hello-svc
spec:
  ports:
  - port: 8080
  selector:
    app: hello-world    <<==== Send to Pods with these labels
    env: tkb            <<==== Send to Pods with these labels
```

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-deploy
spec:
  replicas: 10
  <Snip>
  template:
    metadata:
      labels:
        app: hello-world    <<==== Pod labels
        env: tkb            <<==== Pod labels
    spec:
      containers:
```


EndpointSlice object

- created by the Service object
- Kubernetes curates EndpointSlice lists
 - constantly watch Services with their selectors and Pods with their labels
 - all healthy Pods matching selectors are added to the Service's EndpointSlice object
- When we access the Service IP, our request gets routed through the IP's from the EndpointSlice
- You can query KubeAPI and directly get EndpointSlice content
 - in a Microservice worlds, this is called Service Discovery

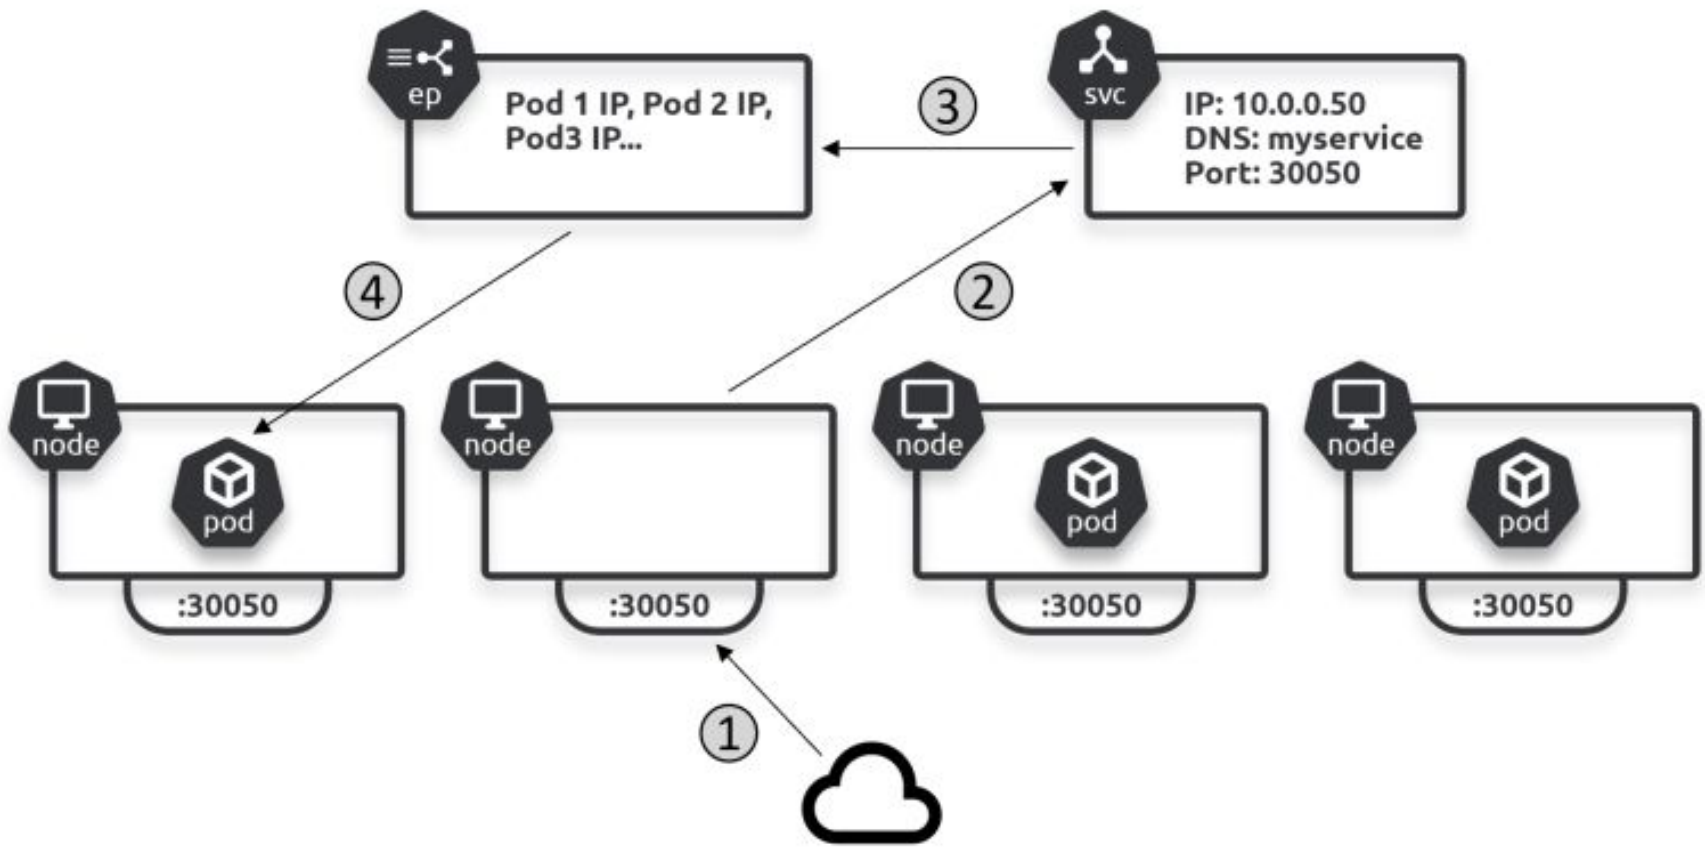
Access a Service from inside

- Service type: ClusterIP
 - our Service gets an IP from the internal K8S network overlay
 - this means our Service can be accessed only from inside
- DNS
 - Services get assigned A and AAAA records:
 - my-svc.my-namespace.svc.cluster-domain.example
 - <https://kubernetes.io/docs/concepts/services-networking/dns-pod-service/#services>
- Stable
 - ClusterIP type will generate
 - stable IP
 - stable DNS

Access a Service from the outside

- We can access Kubernetes objects from the outside in two ways:
 - NodePort
 - LoadBalancer
- NodePort builds on top of ClusterIP, opening the same port on **every** worker node and routing this port to the designated Service
 - You're guessing right, only one deployed Service can be bound to a given port
 - Not scalable
 - Not usable on production clusters
 - with some exceptions, if you know what you're doing

NodePort access



Before we create Services

- There are several ways to get access to the cluster from the inside
 - Ingress (later on)
 - NodePort (now)
- Let's create K3D with NodePort exposed
 - destroy old one
 - `k3d cluster delete kube3`
 - create new one
 - `k3d cluster create kube4 -p "30001:30001@agent:0" -s 1 -a 1`
 - Creates port mapping to the worker and inside
 - https://k3d.io/v5.4.9/usage/exposing_services/?h=nodeport#2-via-nodeport

Create a Service the imperative way

- First of all - not the right way
- Go to the service folder
- Deploy a Deployment
 - `kubectl apply -f deploy.yml`
- Create an imperative Service with NodePort access
 - `kubectl expose deployment svc-test --type=NodePort`
- Get Service info
 - `kubectl describe service svc-test`

Service description

```
$ kubectl describe svc svc-test
```

```
Name:                svc-test
Namespace:           default
Labels:              <none>
Annotations:         <none>
Selector:            chapter=services
Type:               NodePort
IP Family Policy:   SingleStack
IP Families:        IPv4
IP:                 10.43.56.24
IPs:                10.43.56.24
Port:               <unset> 8080/TCP
TargetPort:         8080/TCP
NodePort:           <unset> 30013/TCP
Endpoints:          10.42.0.19:8080,10.42.0.20:8080,10.42.0.21:8080 + 7 more...
Session Affinity:   None
External Traffic Policy: Cluster
Events:             <none>
```



Thoughts

- We can't reach this Service
 - port 30013 was not mapped
- Not the right way (imperative)
- Let's destroy it

Deploy a Service, the declarative way because This is the way

```
apiVersion: v1
kind: Service
metadata:
  name: svc-test
spec:
  type: NodePort
  ports:
    - port: 8080
      nodePort: 30001
      targetPort: 8080
      protocol: TCP
  selector:
    chapter: services
```



Get Services, describe ours and access it

- Get Services

- \$ kubectl get service

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.43.0.1	<none>	443/TCP	15m
svc-test	NodePort	10.43.188.110	<none>	8080:30001/TCP	38s

- Describe our Service

- \$ kubectl describe service svc-test

```
Name: svc-test
Namespace: default
Labels: chapter=services
Annotations: <none>
Selector: chapter=services
Type: NodePort
IP Family Policy: SingleStack
IP Families: IPv4
IP: 10.43.188.110
IPs: 10.43.188.110
Port: <unset> 8080/TCP
TargetPort: 8080/TCP
NodePort: <unset> 30001/TCP
Endpoints: 10.42.0.10:8080,10.42.0.6:8080,10.42.0.7:8080 + 7 more...
Session Affinity: None
External Traffic Policy: Cluster
Events: <none>
```



EndpointSlices

```
$ kubectl get endpointslices
```

NAME	ADDRESSTYPE	PORTS	ENDPOINTS	AGE
svc-test-n7jg4	IPv4	8080	10.42.1.16,10.42.1.17,10.42.0.19 + 7 more...	2m1s
svc-test-9s6sq	IPv6	8080	fd00:10:244:1::c,fd00:10:244:1::9 + 7 more...	2m1s

```
$ kubectl describe endpointslice svc-test-n7jg4
```

Name: svc-test-n7jg4
Namespace: default
Labels: chapter=services
endpointslice.kubernetes.io/managed-by=endpointslice-controller.k8s.io
kubernetes.io/service-name=svc-test

Annotations: endpoints.kubernetes.io/last-change-trigger-time: 2022-01-10T16:20:46Z

AddressType: IPv4

Ports:

Name	Port	Protocol
<unset>	8080	TCP

Endpoints:

- Addresses: 10.42.1.16
Conditions:
Ready: true
Hostname: <unset>
TargetRef: Pod/svc-test-9d7b4cf9d-hnvbf
NodeName: k3d-tkb-agent-2
Zone: <unset>
- Addresses: 10.42.1.17

<Snip>

Events: <none>



LoadBalancer Service type

```
apiVersion: v1
kind: Service
metadata:
  name: cloud-lb
spec:
  type: LoadBalancer
  ports:
    - port: 9000
      targetPort: 8080
  selector:
    chapter: services
```



Ingress

- Ingress exposes HTTP and HTTPS routes from outside the cluster to services within the cluster
- Traffic routing is controlled by rules defined on the Ingress resource
- <https://kubernetes.io/docs/concepts/services-networking/ingress/>



Ingress Controllers

- Mostly used:
 - nginx
 - traefik
- Many many more
 - <https://kubernetes.io/docs/concepts/services-networking/ingress-controllers/>
- You can have more than one Ingress Controller on your cluster
 - control which one is used by a particular ingress object by `spec.ingressClassName` in the ingress manifest

K3D Ingress Controller

- K3D implements K3S in docker
- K3S has traefik as Ingress Controller by default on a Load Balancer node (service of LoadBalancer type)
 - `$ kubectl get deployment traefik --namespace kube-system -o yaml`
 - `$ kubectl get svc traefik --namespace kube-system -o yaml`
- You need to expose port 80 on the LB worker
 - destroy current cluster
 - `k3d cluster delete k3d`
 - create new one with port forwarding to traefik
 - `k3d cluster create kube4 -p "80:80@loadbalancer" -s 1 -a 1`

Ingress Object

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: minimal-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx-example
  rules:
  - http:
      paths:
      - path: /testpath
        pathType: Prefix
        backend:
          service:
            name: test
            port:
              number: 80
```


Deploy example app

- Go to ingress folder
- `kubectl apply -f app.yml`

Simple traefik path example

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: mcu-all
  annotations:
    ingress.kubernetes.io/ssl-redirect: "false"
spec:
  rules:
  - http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: svc-shield
            port:
              number: 8080
```

Complex traefik example

- traefik does not support rewrite target
 - it needs special middleware for these rules
- see k3ing-complex.yml file from the course

Complex traefik - middleware

```
apiVersion: traefik.containo.us/v1alpha1
kind: Middleware
metadata:
  name: mcu-all
spec:
  stripPrefix:
    forceSlash: false
    prefixes:
      - /shield
      - /hydra
  ---
```



Complex traefik - path example

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: mcu-all
  annotations:
    ingress.kubernetes.io/ssl-redirect: "false"
    traefik.ingress.kubernetes.io/router.middlewares: default-mcu-all@kubernetescrd
spec:
  rules:
    - http:
        paths:
          - path: /shield
            pathType: Prefix
            backend:
              service:
                name: svc-shield
                port:
                  number: 8080
          - path: /hydra
            pathType: Prefix
            backend:
              service:
                name: svc-hydra
                port:
                  number: 8080
```

Complex traefik - host example

```
- host: shield.mcu.com
  http:
    paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: svc-shield
            port:
              number: 8080
- host: hydra.mcu.com
  http:
    paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: svc-hydra
            port:
              number: 8080
```

- don't forget to edit your hosts file:
 - 127.0.0.1 shield.mcu.com hydra.mcu.com

Inspecting ingress objects

```
$ kubectl get ing
```

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
mcu-all	traefik	shield.mcu.com,hydra.mcu.com	192.168.64.2,192.168.64.3	80	9m44s

Describe an ingress object

```
Name:          mcu-all
Labels:        <none>
Namespace:     default
Address:       192.168.64.2,192.168.64.3
Ingress Class: traefik
Default backend: <default>
Rules:
  Host          Path  Backends
  ----          -
  *
                /shield svc-shield:8080 (10.42.1.6:8080)
                /hydra   svc-hydra:8080 (10.42.1.5:8080)
  shield.mcu.com
  hydra.mcu.com  /    svc-shield:8080 (10.42.1.6:8080)
                /    svc-hydra:8080 (10.42.1.5:8080)
Annotations:    ingress.kubernetes.io/ssl-redirect: false
                traefik.ingress.kubernetes.io/router.middlewares: default-mcu-all@kubernetescrd
Events:         <none>
```


Clean up your cluster

- k3d cluster delete kube4



Questions?



Reading materials

- The Kubernetes Book (2023 edition)
 - Nigel Poulton
 - Pushkar Joglekar
- Chapter 7: Kubernetes Services
- Chapter 8: Ingress

Class Assignment / Homework

- Create deployments and services for two Pods:
 - One for the app version 1.0 from the last homework
 - One for the app version 2.0 from the last homework
- Create ingress pointing to the two apps:
 - path based: localhost/ver1
 - path based: localhost/ver2
 - host based: ver1.<index>.com
 - host based: ver2.<index>.com
- Deploy manifests and ingress
- Access the four ingress rules in your local browser

